

P-015 | AI FRAUD INVESTIGATION COPILOT

Pricing, GTM and evidence memo | Nguyen Quang Huy | 2A202601873

WHAT IS SOLD		A grounded investigation package for each eligible suspicious alert, embedded into an existing fraud-operations workflow. The AI assembles evidence and policy context; an authorized human owns the final disposition.			
7,648	22,945	32,000 + 30M	82.8%	80% EVAL	44.3%
VND / JOB	PRICE FLOOR	USAGE + PLATFORM	BLENDED GM	CONTAINMENT PROXY	GM BREAK-EVEN
01 PRICING <i>Sell the software workflow, not an unproven autonomous outcome</i> Customer: Mid-market fintech / payment processor Buyer / budget: Head of Risk or COO / software and fraud-operations tooling One job: Complete when the copilot returns a schema-valid, grounded investigation package with required evidence/decision fields; customer final disposition is recorded separately. Value metric: Hybrid: 30M VND platform fee/month + 32,000 VND per completed job. Base fee: 30M/month purchases account-level workflow access, integration/configuration and governance/auditability; it is not a second charge for the package itself. Unit GM: 76.1% on variable usage; 82.8% blended after the 30M/month platform fee at the base case. Rationale: Attribution 2/5: the 5-case local control eval produced 4/5 grounded packages (80%), but no customer time-saved or outcome measurement. Autonomy 1/5: only 1/5 completed without required human intervention; 4/5 cases route to review. Hybrid follows the low-attribution/low-autonomy matrix and avoids unsupported outcome billing. Value anchor: 15 min at 250,000 VND/hour = 62,500 VND/job full value; 50% capture = 31,250 VND/job. Proposed 32,000 VND is above the 22,945 VND floor and below full labor value. RED TEAM: if the 80% package-completion proxy is wrong by 2x to 40%, usage GM falls to 56.1%. Production completion is the economic gate.		02 GTM <i>One channel for 90 days: Sales-Led</i> Affordability: ARPU 106.8M VND/month ACV 1.282B VND / \$49.3K reference 12-month payback CAC math: Budget 1.061B estimated CAC 350M gap 711M affordability ratio 3.03x Deals / AE: 2.81 deals/year 0.013 deals/selling day; mathematically feasible but planning-only. Pain Moment: 09:00-11:00 after an overnight alert burst: a fraud analyst triages high-risk transactions in the existing fraud-operations alert queue/dashboard and needs a grounded case package before manual disposition. Surface: User-facing: existing fraud-analyst alert investigation queue/dashboard side panel. Backend: Kafka fraud_alerts -> fraud engine/agent -> REST /api/v1/fraud/analyze. Live integration is PARTIAL / NOT VERIFIED. Month 1 Learn: 2 design partners, 8 interviews, 300 labeled jobs; produce pain notes, failure taxonomy and time-on-task baseline. Months 2-3 Leverage: 2 pilots / 6,000 jobs / >=85% completion / >=70% GM / CAC <=350M; track 12 qualified opportunities and 3 wins. Month 4+ Expand: Only after >=85% completion, 2 paying customers and security-gap closure. Next niche: multi-merchant payment processors.		03 EVIDENCE <i>Procurement can see what is proven and what is open</i> Eval Results: PARTIAL. Local control eval: 5 attempted, 4 grounded packages (80% commercial completion proxy), 1 autonomous completion (20%), 4 human-review cases and 1 grounding failure; retry telemetry is not instrumented. Eight RAG queries: 100% citation accuracy. Upstream ML recall: 80.08% validation / 84.84% integration. Not customer evidence. Risk Checklist: PARTIAL. Grounding, fallback, injection defense, score preservation, redaction and human review are evidenced. Auth/RBAC, rate limit, TLS, retention/deletion, durable audit export, vendor terms and live deployment remain open. Pilot Report: MISSING EVIDENCE. No customer or participant is fabricated. PLANNED - NOT YET EXECUTED: start 2026-10-01, end 2026-11-26, 2 design partners, 6,000 jobs; owner Nguyen Quang Huy; measure completion, time saved, latency, retries, Cost/Job and safety. Submission gate: READY WITH DISCLOSED LIMITATIONS. Official Day28 templates were not found locally; this memo and workbook are transparent substitutes. Numeric claims trace to workbook cells in evidence/onepager-traceability.md. Production readiness remains gated. Decision: READY for artifact review/submission. Production claims remain gated on the planned pilot and open controls.	

Source trail: P-015 artifacts/agent + artifacts/ml/lightgbm; P-015 docs/business_product_direction.md, ai_fraud_agent.md, security.md, deployment_guide.md; current OpenAI, Supabase, Stripe Radar and Fingerprint pricing checked 2026-08-27. See workbook and evidence pack for full notes.